

ALI NIKKHAH

Open to Relocation & Remote | alinkkh9@gmail.com | +98 991 296 3951 | alinikkhah2001.github.io |
linkedin.com/in/alinikkhah2001

SUMMARY

Applied ML researcher & engineer (Sharif B.Sc. Electrical Engineering 2020-2025, GPA 3.65; M.Sc. AI from Sep 2025) specialising in agentic RAG, multimodal vision-language, self-hosted LLM serving, and GPU runtime optimisation. Shipped LLM + CV pipelines at marketplace scale, self-hosted Graph RAG assistant for a national credit bureau, and BNPL risk platforms. Research focus: LLM reasoning and hallucination mitigation. — Variant: ML Platform

TECHNICAL SKILLS

Languages: Python, C++, Go, SQL, R, Bash, TypeScript
ML/Vision: PyTorch, TensorFlow, JAX, Hugging Face, ViT, BLIP, T5, Whisper
Agentic: LangGraph, LangChain, LlamaIndex, FAISS, ChromaDB, HyDE
LLM Serving/GPU: vLLM, TGI, AWQ, GPTQ, KV-cache
Data/MLOps: Spark, Airflow, Kafka, Docker, Kubernetes, Ray

PROFESSIONAL EXPERIENCE

Iran Credit Scoring Bureau

Tehran, Iran

ML Engineer — Self-Hosted Graph RAG Assistant (Concurrent)

Apr 2026 == Present

- Architected self-hosted Graph RAG assistant: Neo4j knowledge graph + vector retrieval fused into an agent that answers analyst queries over credit-scoring domain data with grounded citations.
- Ran self-hosted open-weight LLMs (vLLM/TGI) on GPU; profiled and optimised GPU runtime — KV-cache, batching, quantization (AWQ/GPTQ) — cutting per-token latency and GPU memory footprint.
- Implemented Neo4j tool-calls (Cypher generation + schema-aware retrieval) so the assistant executes graph traversals instead of hallucinating structured facts.
- Added hallucination guardrails — retrieval grounding, citation enforcement, deterministic fallback — for a regulated financial environment.
- *Stack:* Python, Neo4j, Cypher, Graph RAG, LangGraph, LangChain, FAISS, vLLM
- *Most recent Iran position; overlaps Sharif M.Sc. AI (Sep 2025) — full-time industry alongside full-time graduate study*

Turquoise Digital

Tehran, Iran

Senior Data Scientist & Data Engineer (Concurrent)

Sep 2025 == Mar 2026

- Centralized lake from raw CDC via PeerDB + Trino/Parquet on 5TB/day, <60s freshness; Airflow + PySpark DAGs cut memory 60%, governed by Great Expectations + Prometheus/Grafana.
- Built default-risk ensemble (XGBoost + time-series Transformers) — AUC 0.87, -22% defaults; 50+ concurrent mSPRT A/B tests with canary & auto-rollback, presented to C-suite.
- Shipped Whisper/wav2vec2.0 + LangGraph agentic chatbot — -78% ticket volume, 94% intent accuracy; deterministic fallback + policy guardrails for auditability.
- Drove portfolio allocation & risk engine (Sharpe +0.42) with streaming ClickHouse signals; LlamaIndex retrieval graph for ranking features.
- *Stack:* Python, PyTorch, Spark, Trino, PeerDB, Airflow, Great Expectations, Prometheus
- *Concurrent with Advanced Analytics Australia (Remote Contract) — 7-month overlap, split focus: Turquoise BNPL/risk/whisper vs Advanced edge CV*

Advanced Analytics Australia

Remote · Australia (Tehran-based)

Computer Vision / ML Platform Engineer (Edge) (Concurrent)

Sep 2025 == Apr 2026

- Deployed YOLO + ViT quantized (TensorRT/TF Lite) on Triton — 99.2% uptime <50ms p99, 3.5x speedup, exactly-once Kafka, 30 FPS over 100+ streams.
- Built edge MLOps with Triton model repository, zero-downtime rolling updates, automated rollback, Prometheus/Grafana per-site health.
- Cut false positives -35% via attention-guided NMS + SNR monitoring; broker architecture for 12 distributed sites.
- *Stack:* Python, OpenCV, YOLO, ViT, TensorRT, TFLite, Triton, Kafka
- *Parallel remote contract alongside Turquoise — edge CV track*

Digikala

Tehran, Iran

AI Engineering R&D Lead (Concurrent)

Jan 2025 == Aug 2025

- Led team of 5 defining R&D roadmap + SLOs; shipped MCP with LangGraph cyclic graphs, zero-downtime hot-swapping, health-checked agent composition with deterministic fallback.
- Hybrid retrieval FAISS + Elasticsearch RRF + HyDE + KAG — +23% recall@10, -41% hallucination; multi-tier Redis caching, p99 <200ms.

- LLM serving on vLLM (PagedAttention) + Triton — 3.2x throughput, 60% cost reduction; 200 pods, MIG/HPA/VPA, Prometheus KV-cache dashboards.
- Partnered with infra/product establishing eval pipelines, A/B guardrails, canary rollouts.
- *Stack*: Python, LangGraph, LangChain, LlamaIndex, FAISS, Elasticsearch, ChromaDB, HyDE
- *Full-time lead; UBC/Trinity/Sharif were part-time remote (<10h/wk)*

Digikala

Tehran, Iran

Nov 2024 == Jan 2025

AI Engineer (Concurrent)

- Implemented LangGraph orchestration layers with health monitoring; hybrid retrieval with RRF + deterministic fallbacks for auditability.
- Embedded eval telemetry, memory-efficient indexing, response caching; load-tested to 5M queries.
- *Stack*: LangChain, LlamaIndex, FAISS, Elasticsearch, FastAPI
- *Promotion track — converted to R&D Lead Jan 2025*

HomaCloud

Tehran, Iran

Aug 2021 == Feb 2022

MLOps Engineer (Founding) (Concurrent)

- Provisioned multi-node K8s via Terraform; Ray + MLflow training pipelines, TF Serving + FastAPI — <10 min releases, 99.9% deploy success, 15+ models.
- GitOps with GitHub Actions, canary, spot-instance optimization; Prometheus/Grafana + ELK observability; registry + lineage.
- *Stack*: Kubernetes, Docker, Ray, MLflow, Terraform, ArgoCD, FastAPI, Prometheus
- *During B.Sc. — clean sequential after Mapna internship*

SELECTED PROJECTS & RESEARCH

Trinity College Dublin · EmoDub <i>Emotion-aware voice translation</i> - SER front-end + wav2vec2.0 bullets (see experiences)	May 2025 – Nov 2025
University of British Columbia <i>Ultrasound report generation</i>	Feb 2024 – Nov 2025
L3S Research Center <i>Texture-free motion intelligence</i> Automated DICOM-Compliant Ultrasound Report Generation via Contrastive Vision-Language Learning — <i>arXiv preprint</i> — <i>in preparation</i> , 2025	Apr 2023 – Feb 2024

EDUCATION

Sharif University of Technology <i>M.Sc. Artificial Intelligence Computer Engineering (full-time, ongoing)</i>	Tehran, Iran Sept 2025 – Present
Sharif University of Technology <i>B.Sc. Electrical Engineering GPA: 3.65</i>	Tehran, Iran Sept 2020 – Sept 2025