

Ali Nikkhah

Machine Learning Engineer / Researcher

📍 Tehran, Iran · Remote-friendly · Open to Relocation ✉ alinkkh9@gmail.com ☎ 0991 296 3951
🔗 alinikkhah2001.github.io 🌐 alinikkhah2001 🔄 AliNikkhah2001

Summary

Applied ML researcher & engineer (Sharif B.Sc. Electrical Engineering 2020-2025, GPA 3.65; M.Sc. AI from Sep 2025) specialising in agentic RAG, multimodal vision-language, self-hosted LLM serving, and GPU runtime optimisation. Shipped LLM + CV pipelines at marketplace scale, self-hosted Graph RAG assistant for a national credit bureau, and BNPL risk platforms. Research focus: LLM reasoning and hallucination mitigation.

Education

M.Sc. Artificial Intelligence	Sharif University of Technology , M.Sc. Artificial Intelligence	Tehran, Iran
	• Computer Engineering Department, AI major (full-time, ongoing, no thesis yet)	Sept 2025 – present
B.Sc. Electrical Engineering	Sharif University of Technology , B.Sc. Electrical Engineering	Tehran, Iran
	• GPA: 3.65	Sept 2020 – Sept 2025
	• Thesis: Team thesis on efficient design and implementation of spread-time CDMA ASIC (my slice: FPGA implementation and pre-chip FPGA tests), supervisor Dr. Fatemeh (Noyan) Akbar, grade 20/20	
	• Coursework: Machine Learning, Deep Learning, Natural Language Processing, Data Science, Embedded Systems, Signals and Systems, Microprocessor Systems	

Experience

Turquoise Digital , Senior Data Scientist & Data Engineer	Tehran, Iran — Full-time · On-site
Directed AI for regional retailers — multilingual reco, BNPL credit-risk ensemble (AUC 0.87), agentic support chatbot, centralized data lake + observability.	Sept 2025 – Mar 2026
• Centralized lake from raw CDC via PeerDB + Trino/Parquet on 5TB/day, <60s freshness; Airflow + PySpark DAGs cut memory 60%, governed by Great Expectations + Prometheus/Grafana.	7 months
• Built default-risk ensemble (XGBoost + time-series Transformers) — AUC 0.87, -22% defaults; 50+ concurrent mSPRT A/B tests with canary & auto-rollback, presented to C-suite.	
• Shipped Whisper/wav2vec2.0 + LangGraph agentic chatbot — -78% ticket volume, 94% intent accuracy; deterministic fallback + policy guardrails for auditability.	
• Drove portfolio allocation & risk engine (Sharpe +0.42) with streaming ClickHouse signals; LlamaIndex retrieval graph for ranking features.	
• Stack: Python, PyTorch, Spark, Trino, PeerDB, Airflow	
• Concurrent with Advanced Analytics Australia (Remote Contract) — 7-month overlap, split focus: Turquoise BNPL/risk/whisper vs Advanced edge CV	

Advanced Analytics Australia , Computer Vision / ML Platform Engineer (Edge) Edge CV platform across 12 retail sites, 100+ cameras @30 FPS, Kafka streaming, quantized TensorRT/TF Lite + Triton serving. • Deployed YOLO + ViT quantized (TensorRT/TF Lite) on Triton — 99.2% uptime <50ms p99, 3.5x speedup, exactly-once Kafka, 30 FPS over 100+ streams. • Built edge MLOps with Triton model repository, zero-downtime rolling updates, automated rollback, Prometheus/Grafana per-site health. • Cut false positives -35% via attention-guided NMS + SNR monitoring; broker architecture for 12 distributed sites. • Stack: Python, OpenCV, YOLO, ViT, TensorRT, TFLite • Parallel remote contract alongside Turquoise — edge CV track	Remote · Australia (Tehran-based) — Remote Contract Sept 2025 – Apr 2026 8 months
Digikala , AI Engineering R&D Lead Owned agentic RAG platform for marketplace listing intelligence — modular control pipelines (MCP), hybrid retrieval, LLM serving at 5M+ queries/day. • Led team of 5 defining R&D roadmap + SLOs; shipped MCP with LangGraph cyclic graphs, zero-downtime hot-swapping, health-checked agent composition with deterministic fallback. • Hybrid retrieval FAISS + Elasticsearch RRF + HyDE + KAG — +23% recall@10, -41% hallucination; multi-tier Redis caching, p99 <200ms. • LLM serving on vLLM (PagedAttention) + Triton — 3.2x throughput, 60% cost reduction; 200 pods, MIG/HPA/VPA, Prometheus KV-cache dashboards. • Partnered with infra/product establishing eval pipelines, A/B guardrails, canary rollouts. • Stack: Python, LangGraph, LangChain, LlamaIndex, FAISS, Elasticsearch • Full-time lead; UBC/Trinity/Sharif were part-time remote (<10h/wk)	Tehran, Iran — Full-time · On-site Jan 2025 – Aug 2025 8 months
HomaCloud , MLOps Engineer (Founding) Built K8s-native training, CI/CD, and serving infra from scratch — Docker/K8s, Ray, MLflow, FastAPI/TF Serving, Terraform, ArgoCD. • Provisioned multi-node K8s via Terraform; Ray + MLflow training pipelines, TF Serving + FastAPI — <10 min releases, 99.9% deploy success, 15+ models. • GitOps with GitHub Actions, canary, spot-instance optimization; Prometheus/Grafana + ELK observability; registry + lineage. • Stack: Kubernetes, Docker, Ray, MLflow, Terraform, ArgoCD • During B.Sc. — clean sequential after Mapna internship	Tehran, Iran — Full-time · On-site Aug 2021 – Feb 2022 7 months

Publications

Automated DICOM-Compliant Ultrasound Report Generation via Contrastive Vision-Language Learning Ali Nikkhah, [Co-authors TBD], A. Supervisor (arXiv preprint — in preparation)	Jan 2025
---	----------

Teaching

Natural Language Processing (Graduate)

TA · Sharif (Aug 2023 – Jan 2024)

Generative Models (Graduate)

TA · Sharif (Aug 2023 – Jan 2024)

Deep Learning (Graduate)

TA · Sharif (Jan 2023 – Aug 2023)

Analog Electronics

TA · Sharif (Aug 2022 – Jan 2023)

Machine Learning & Data Science

TA · Sharif (Jan 2022 – Aug 2022)

Circuit Theory & Engineering Probability

TA · Sharif (Jan 2021 – Aug 2021)

Skills

Programming: Python, C++, Go, SQL, R, Bash, TypeScript

ML & Vision: PyTorch, TensorFlow, JAX, Hugging Face, ViT, BLIP, T5, Whisper, wav2vec2.0, Librosa, OpenCV, YOLO, ONNX

Agentic AI & RAG: LangGraph, LangChain, LlamaIndex, FAISS, ChromaDB, HyDE, KAG, Graph RAG, vLLM, Triton, Neo4j, Cypher

LLM Serving & GPU: vLLM, TGI, AWQ, GPTQ, KV-cache, GPU Optimization, PagedAttention, TensorRT-LLM

Data & MLOps: Spark, Airflow, Kafka, Docker, Kubernetes, Ray, MLflow, W&B, ArgoCD, Terraform

Languages: Persian (Native), English (C2 IELTS 8.0), Arabic (Professional), French (Limited), German (Limited)